

Task 2 — Values Audit

RegBridge Analytics and the Design Values Behind the Tool

2.1 Company Mission and Identity

For this assignment, I assume that the tool is developed by a hypothetical RegTech company called RegBridge Analytics. I define the company’s mission as follows:

RegBridge Analytics helps financial institutions use AI credit decisioning tools responsibly across jurisdictions by connecting regulatory rules, model explainability, fairness monitoring, and human review.

I see RegBridge Analytics as a small but specialised B2B RegTech company, not a general compliance software provider. At this stage, it is an early-stage company with around 25 employees, including regulatory analysts, machine learning engineers, product designers, and financial risk specialists. Its target clients are banks, credit card issuers, BNPL platforms, and consumer finance companies that use algorithmic systems in credit decisions.

The company’s first market focus is the United States and the European Union, because these two jurisdictions ask different questions about AI credit decisioning. In the US, the concern is strongly linked to fair lending and adverse action explanations. The CFPB has made clear that lenders using AI or complex models still need to give specific and accurate reasons when denying credit (Consumer Financial Protection Bureau [CFPB], 2023). In the EU, creditworthiness assessment and credit scoring systems are treated as high-risk AI systems under the EU AI Act, which brings broader expectations around governance, risk management, data quality, transparency, and human oversight (European Union, 2024). The NIST AI Risk Management Framework also supports the idea that AI risks should be governed, measured, and managed throughout the system lifecycle, rather than only documented after problems appear (National Institute of Standards and Technology [NIST], 2023).

I would describe the company’s values in four practical terms:

Value	What it means for my tool
Substance over box-ticking	The tool should identify real model and consumer harm risks, not just produce clean compliance files.
Explainability	If a consumer is denied credit, the reasons should be specific and understandable, not hidden behind a “model score.”
Human accountability	The tool should support judgement, not replace it. Human reviewers remain responsible for difficult cases.
Jurisdictional awareness	The tool should not pretend that one compliance logic works equally well in every jurisdiction.

These values matter because this assignment is not only about building a technical prototype. It is also about recognising that RegTech tools reflect choices about what risks matter, who should be protected, and how

much judgement should remain with humans.

2.2 Whose Perspective Does the Tool Serve?

The direct client of FairCredit Jurisdiction Navigator would be the compliance, model risk, or credit risk function of American Express. This is realistic because these teams would face the most immediate pressure if the company could not explain credit decisions, detect unfair outcomes, or show proper governance over AI-supported credit approval.

However, I do not think the tool should serve only the paying client. In my design, there are three main stakeholder perspectives:

Stakeholder	What they want	Possible tension
American Express compliance team	Clear audit trail, lower regulatory risk, usable workflow	May prefer documentation that is clean and defensible
Regulators	Evidence that AI systems are explainable, fair, and properly governed	May expect the tool to reveal uncomfortable risks
Consumers	Fair treatment and meaningful explanations	Usually have no access to the model, data, or dashboard

The main tension is that the buyer and the most affected stakeholder are not the same. American Express would pay for the tool, but consumers bear much of the harm if the model is unfair or poorly explained. This matters because a tool designed only for the client could become a documentation engine. It might help the firm appear compliant without actually showing whether the credit model is producing harmful outcomes.

My own position is that the tool should serve the client operationally, but not only the client normatively. In practice, the dashboard has to be useful for American Express, otherwise it will not be adopted. But the internal logic of the tool should still be anchored in genuine risk detection. That means the tool should sometimes produce outputs that are inconvenient for the client, such as a bias alert, model drift warning, or mandatory human review flag.

This is also why I chose AI credit decisioning as the domain. A credit decision is not only a private commercial decision. It affects access to financial opportunity. In the US, this appears through the requirement to give specific adverse action reasons. In the EU, it appears through the classification of credit scoring systems as high-risk AI. Both regimes show that credit decisions have social consequences beyond ordinary business efficiency (CFPB, 2023; European Union, 2024).

2.3 Is the Tool Measuring Genuine Risk or Documenting Compliance?

My intention is for the tool to measure genuine risk, although it will also produce documentation as a by-product. This distinction is important. A weak version of this tool could simply generate an adverse action template for the US and an AI Act checklist for the EU. That would be useful for filing purposes, but it would

not necessarily show whether the model is causing real harm.

To avoid that, I include one design choice that reflects the values of the company: the EU mode does not only generate a governance checklist. It also calculates portfolio-level risk indicators, including: approval-rate disparity between applicant groups; false-positive and false-negative rates by group; model performance drift; population stability index; concentration of adverse action reasons across groups.

These metrics are important because documentation can look complete while the model is still behaving unfairly. For example, a model can have a model card, a review date, and an audit trail, but still approve one group of applicants at a much lower rate than another group. Similarly, a model may perform well overall while being weak for a smaller subgroup. In these cases, the real problem is not missing paperwork; it is that the model may be producing unequal or unstable outcomes.

Another design choice is that protected attributes or proxy group indicators are not used to train the credit model in the prototype. Instead, they are used only for post-model fairness monitoring. I chose this because I do not want the model to directly learn from sensitive group identity. At the same time, I do not want to pretend that removing protected attributes automatically removes discrimination. Variables such as region, employment pattern, credit history length, or income stability can still act as proxies. Therefore, the tool needs a fairness-monitoring layer after the model has produced its decisions.

This is not a perfect solution, but I think it is a more honest one. It recognises that algorithmic fairness cannot be solved only by deleting sensitive columns. The harder task is to test whether the final outcomes are still unequal, unstable, or difficult to justify. This approach is consistent with the broader risk-management logic of the NIST AI RMF, which emphasises structured governance and risk measurement rather than treating AI risk as a one-time compliance exercise (NIST, 2023).

2.4 Who Bears the Cost If the Tool Gets It Wrong?

The cost of failure depends on the type of mistake. I summarise the main risks below:

Failure type	What could happen	Main stakeholder harmed
False positive	A fair or compliant model is wrongly flagged as risky	American Express, because of review cost and operational delay
False negative	An unfair or poorly explained model is not flagged	Consumers, because they may be denied credit unfairly
Model drift	The model becomes less reliable as applicant behaviour or economic conditions change	Both consumers and American Express
Jurisdictional misconfiguration	EU cases are assessed using US-style rules only, or vice versa	American Express, regulators, and consumers

Failure type	What could happen	Main stakeholder harmed
Over-trust in the tool	Staff treat the tool output as a final answer	All stakeholders, because human judgement is weakened

Among these, I think false negatives are the most serious. If the tool fails to detect unfair outcomes or vague adverse action explanations, the harm is not just administrative. Consumers may lose access to credit without a clear reason, and the firm may not realise that its model is producing biased or unstable results.

Jurisdictional misconfiguration is also a central risk for this project. For example, if an EU credit portfolio is reviewed only through the US adverse action lens, the tool may still produce individual explanations but fail to trigger EU-style high-risk AI governance checks. In that case, American Express might believe it has addressed the compliance issue, when it has actually applied the wrong regulatory logic.

This is why the tool should include clear human review boundaries. It should not make the final legal or ethical judgement. It should classify, explain, flag, and document. Final decisions should remain with compliance officers, model risk specialists, and legal teams, especially when the issue involves fairness, jurisdictional interpretation, or borderline credit outcomes.

2.5 Final Values Position

My values position is that the tool should be commercially useful, but not morally neutral. RegBridge Analytics is building a product for a paying financial institution, so the tool must fit the client's workflow and reduce real operational burden. However, its legitimacy depends on whether it helps identify real risks to consumers and regulators, not only whether it creates good-looking compliance evidence.

For that reason, I design FairCredit Jurisdiction Navigator around three principles:

1. Explain individual decisions clearly where adverse action explanation is central. This is especially important in the US mode.
2. Monitor system-level fairness and governance where high-risk AI regulation is central. This is especially important in the EU mode.
3. Route uncertainty to human judgement instead of hiding it behind automated confidence scores. This applies to both jurisdictions.

In simple terms, the US mode asks: Can American Express explain this specific negative credit decision?

The EU mode asks: Is this AI credit system being governed and monitored properly as a high-risk system?

That difference is the reason the tool needs to be jurisdiction-aware. The tool should not assume that the US and EU are asking the same question. It should make those differences visible in the design of the dashboard, the metrics, the alerts, and the human review process.

References List:

1. Consumer Financial Protection Bureau. (2023, September 19). CFPB issues guidance on credit denials by lenders using artificial intelligence.
<https://www.consumerfinance.gov/about-us/newsroom/cfpb-issues-guidance-on-credit-denials-by-lenders-using-artificial-intelligence/>
2. European Union. (2024). Regulation (EU) 2024/1689 of the European Parliament and of the Council laying down harmonised rules on artificial intelligence. Official Journal of the European Union.
https://eur-lex.europa.eu/legal-content/EN/TXT/PDF/?uri=OJ:L_202401689
3. National Institute of Standards and Technology. (2023). Artificial intelligence risk management framework (AI RMF 1.0). U.S. Department of Commerce. <https://www.nist.gov/itl/ai-risk-management-framework>